

```
· using Pkg; Pkg.activate(".")
```

Metalhead Example

```
base_path = "data"
```

```
· base_path = "data"
```

```
· using Metalhead, Flux
```

We will use an image of Philip, the corgi to test out whether our trained ML models perform on new images that they haven't seen before. In other words, do they generalize sufficiently?

```
philip =
```



```
· philip = load(joinpath(base_path, "philip_crop.jpg"))
```

Load the pretrained VGG model

This is a classic image recognition model which has been trained on the ImageNet dataset.

```
vgg = VGG19()
```

```
· vgg = VGG19()
```

```
"Cardigan, Cardigan Welsh corgi"
```

```
· classify(vgg, philip)
```

Excellent! Our model was able to identify that the picture is of Philip, the corgi.

Although, it did also confuse it for a cardigan, but could you really blame it? Philip is one cuddly dog.



Flux3D Example

```
· using Flux3D, Zygote, AbstractPlotting, Statistics
```

```
3×2562 Array{Float32,2}:
0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  ...  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0
0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0      0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0
0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0      0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0

· begin
·     dolphin = load_trimesh(joinpath(base_path, "dolphin.obj"))
·     src = load_trimesh(joinpath(base_path, "sphere.obj"))
·     _offset = zeros(Float32, size(get_verts_packed(src))...)
· end
```

Visualizing Initial Guess and Target Meshes



```
· [load(joinpath(base_path, "src.png")); load(joinpath(base_path, "dolphin.png"))]
```

Normalizing the Target Mesh

```
3×2562 Array{Float32,2}:
-0.101319335  0.100773014  -0.102468416  ...  -0.059885267  -0.061462224  -0.07344546
 0.28716984  0.28668007  -0.15838437  ...  -0.21195507  -0.15824707  -0.17916189
 0.09061999  0.09080836  -0.12292277  ...  -0.64931875  -0.61420053  -0.60773396
```

```
· begin
·     tgt = deepcopy(dolphin)
·     verts = get_verts_packed(tgt)
·     center = mean(verts, dims=2)
·     verts = verts .- center
```

```

·     scale = maximum(abs.(verts))
·     verts = verts ./ scale
·     tgt._verts_packed = verts
· end

```

Defining the Loss

As discussed earlier, we have to define a metric we wish to optimise for - our "loss function".

This would help us demonstrate whether we can effectively learn the dolphin mesh using the differentiable programming primitives.

loss_dolphin (generic function with 1 method)

```

· function loss_dolphin(x::AbstractArray, src::TriMesh, tgt::TriMesh)
·     src = Flux3D.offset(src, x)
·     loss1 = chamfer_distance(src, tgt, 5000)
·     loss2 = laplacian_loss(src)
·     loss3 = edge_loss(src)
·     return loss1 + 0.1*loss2 + loss3
· end

```

Training the Model

To run the actual training loop, set *nepochs* to 2000.

nepochs =

10

```

· nepochs = 10




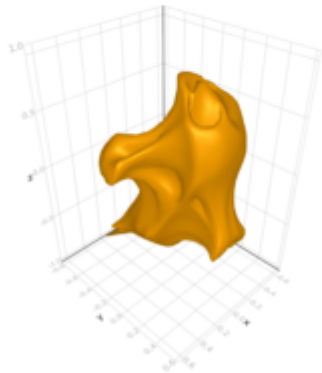

```

```

· begin
·     θ = Flux.params(_offset)
·     lr = 1.
·     opt2 = Momentum(lr, 0.9)
·     for itr in 1:nepochs
·         gs = gradient(θ) do
·             loss_dolphin(_offset, src, tgt)
·         end
·         Flux.update!(opt2, _offset, gs[_offset])
·     end
· end

```

Visualizing the training results



```
[
  load(joinpath(base_path, "src.png"));
```

```

· load(joinpath(base_path, "src_50.png"));
· load(joinpath(base_path, "src_200.png"));
· load(joinpath(base_path, "src_500.png"));
· load(joinpath(base_path, "src_2000.png"));
· ]

```

SciML Example - Hooke's Law

Imperfect Spring


$$x'' = -k \times x + 0.1 \times \sin(x)$$

```

· using DifferentialEquations, Plots

```

Plot Actual Solution




```

· begin
·   k = 1.0
·   force(dx,x,k,t) = -k*x + 0.1sin(x)
·   prob = SecondOrderODEProblem(force,1.0,0.0,(0.0,10.0),k)
·   sol = solve(prob)
·   Plots.plot(sol,label=["Velocity" "Position"])
· end

```

Sampling Datapoints from Actual Solution



```

• begin
•   t = 0:0.001:1.0
•   plot_t = 0:0.01:10
•   data_plot = sol(plot_t)
•   positions_plot = [state[2] for state in data_plot]
•   force_plot = [force(state[1],state[2],k,t) for state in data_plot]
•
•   # Generate the dataset
•   t = 0:3.3:10
•   dataset = sol(t)
•   position_data = [state[2] for state in sol(t)]
•   force_data = [force(state[1],state[2],k,t) for state in sol(t)]
•
•   Plots.plot(plot_t,force_plot,xlabel="t",label="True Force")
•   Plots.scatter!(t,force_data,label="Force Measurements")
• end

```

Define Neural Network and L2 Loss

The neural network is trained to match the force values at every position

loss (generic function with 1 method)

```

• begin
•   NNForce = Chain(x -> [x],
•                   Dense(1,32,tanh),
•                   Dense(32,1),
•                   first)
•
•   loss() = sum(abs2,NNForce(position_data[i]) - force_data[i] for i in
•               1:length(position_data))
• end

```


Train the neural network


```

· begin
·   opt = Flux.Descent(0.01)
·   data = Iterators.repeated((), 5000)
·   Flux.train!(loss, Flux.params(NNForce), data, opt)
· end

```

Plot - Trained Model with 4 Datapoints

The trained model fits the datapoints perfectly, but does not capture the real physics



```


· begin
·   learned_force_plot = NNForce.(positions_plot)
·
·   Plots.plot(plot_t, force_plot, xlabel="t", label="True Force")
·   Plots.plot!(plot_t, learned_force_plot, label="Predicted Force")
·   Plots.scatter!(t, force_data, label="Force Measurements")
· end

```

This shows the neural network fit the force data but doesn't match the force function well enough.

In this case, an extra loss component which combines the constraints of the ideal spring are added

Plot - Ideal Spring vs Real Spring



```

• begin
•   force2(dx,x,k,t) = -k*x
•   prob_simplified = SecondOrderODEProblem(force2,1.0,0.0,(0.0,10.0),k)
•   sol_simplified = solve(prob_simplified)
•   Plots.plot(sol,label=["Velocity - Real" "Position - Real"])
•   Plots.plot!(sol_simplified,label=["Velocity - Ideal" "Position - Ideal"])
• end

```

Generate more data assuming ideal spring. This is done by calculating force positions at random points.

composed_loss (generic function with 1 method)

```

• begin
•   random_positions = [2rand()-1 for i in 1:100] # random values in [-1,1]
•   loss_ode() = sum(abs2,NNForce(x) - (-k*x) for x in random_positions)
•
•   # Composed loss with weighted ODE loss component
•   λ = 0.1
•   composed_loss() = loss() + λ*loss_ode()
• end

```


Train Neural Network with ODE loss component

```

• let
•   opt = Flux.Descent(0.01)
•   data = Iterators.repeated((), 5000)
•   Flux.train!(composed_loss, Flux.params(NNForce), data, opt)
• end

```

Plot - Trained Neural Network vs Actual Force



```
· let  
·   learned_force_plot = NNForce.(positions_plot)  
·  
·   Plots.plot(plot_t, force_plot, xlabel="t", label="True Force")  
·   Plots.plot!(plot_t, learned_force_plot, label="Predicted Force")  
·   Plots.scatter!(t, force_data, label="Force Measurements")  
· end
```

This shows the trained neural network approximating the actual force function very closely.